

Artificial Neural Network

Importing the libraries

In [1]:	<pre>import numpy as np import pandas as pd import tensorflow as tf</pre>
In [2]:	<pre>tf.__version__</pre>
Out[2]:	'2.7.0'

Part 1 - Data Preprocessing

Importing the dataset

In [4]:	<pre>dataset = pd.read_csv('Project_1_ANN_Banking.csv') X = dataset.iloc[:, 3:-1].values #Printing starting from index 3(CreditScore) to EstimatedSalary excluding -1(Exited)--->Train Attributes y = dataset.iloc[:, -1].values #Printing only last column Exited-->Target Attribute</pre>
In [5]:	<pre>print(X)</pre> <pre>[[619 'France' 'Female' ... 1 1 101348.88] [608 'Spain' 'Female' ... 0 1 112542.58] [502 'France' 'Female' ... 1 0 113931.57] ... [709 'France' 'Female' ... 0 1 42085.58] [772 'Germany' 'Male' ... 1 0 92888.52] [792 'France' 'Female' ... 1 0 38190.78]]</pre>
In [6]:	<pre>print(X)</pre> <pre>[[619 'France' 'Female' ... 1 1 101348.88] [608 'Spain' 'Female' ... 0 1 112542.58] [502 'France' 'Female' ... 1 0 113931.57] ... [709 'France' 'Female' ... 0 1 42085.58] [772 'Germany' 'Male' ... 1 0 92888.52] [792 'France' 'Female' ... 1 0 38190.78]]</pre>

Encoding categorical data

In [9]:	Label Encoding the "Gender" column
In [10]:	<pre>from sklearn.preprocessing import LabelEncoder le = LabelEncoder() X[:, 2] = le.fit_transform(X[:, 2]) #Encoding index 2 column only(Gender)</pre>
In [11]:	<pre>print(X)</pre> <pre>[[619 'France' 0 ... 1 1 101348.88] [608 'Spain' 0 ... 0 1 112542.58] [502 'France' 0 ... 1 0 113931.57] ... [709 'France' 0 ... 0 1 42085.58] [772 'Germany' 1 ... 1 0 92888.52] [792 'France' 0 ... 1 0 38190.78]]</pre> <p>One Hot Encoding the "Geography" column</p>
In [13]:	<pre>from sklearn.compose import ColumnTransformer from sklearn.preprocessing import OneHotEncoder ct = ColumnTransformer(transformers=[('encoder', OneHotEncoder(), [1])], remainder='passthrough') #[1]--->Column index 1 X = np.array(ct.fit_transform(X))</pre>
In [14]:	<pre>print(X)</pre> <pre>[[1.0 0.0 0.0 ... 1 1 101348.88] [0.0 0.0 1.0 ... 0 1 112542.58] [1.0 0.0 0.0 ... 1 0 113931.57] ... [1.0 0.0 0.0 ... 0 1 42085.58] [0.0 1.0 0.0 ... 1 0 92888.52] [1.0 0.0 0.0 ... 1 0 38190.78]]</pre>

Splitting the dataset into the Training set and Test set

In [15]:	<pre>from sklearn.model_selection import train_test_split X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 0) #0.2--->20% data from train data set will be tested</pre>
----------	---

Feature Scaling

In [16]:	<pre>from sklearn.preprocessing import StandardScaler sc = StandardScaler() X_train = sc.fit_transform(X_train) X_test = sc.transform(X_test) #Scaling everything</pre>
----------	---

Part 2 - Building the ANN

Initializing the ANN

In [17]:	<pre>ann = tf.keras.models.Sequential() #Initializing ANN as a sequence of layers</pre>
----------	---

Adding the input layer and the first hidden layer

In [18]:	<pre>ann.add(tf.keras.layers.Dense(units=6, activation='relu')) #Automatically add input neurons,assuming 6 neurons in first hidden layer</pre>
----------	---

Adding the second hidden layer

In [19]:	<pre>ann.add(tf.keras.layers.Dense(units=6, activation='relu'))</pre>
----------	---

Adding the output layer

In [21]:	<pre>ann.add(tf.keras.layers.Dense(units=1, activation='sigmoid')) #For probability of binary outcome & prediction sigmoid preferable,Output layer 1</pre>
----------	--

Part 3 - Training the ANN

Compiling the ANN

In [22]:	<pre>ann.compile(optimizer = 'adam', loss = 'binary_crossentropy', metrics = ['accuracy']) #adam-->updating weights to reduce loss through iteration #loss-->compute difference between prediction & real result #binary_crossentropy-->Always be for binary classification & for non-binary it will be catagorical_crossentropy</pre>
----------	---

Training the ANN on the Training set

In [23]:	<pre>ann.fit(X_train, y_train, batch_size = 32, epochs = 100)</pre> <div>Epoch 1/100 250/250 [=====] - 4s 7ms/step - loss: 0.6375 - accuracy: 0.6315 Epoch 2/100 250/250 [=====] - 2s 6ms/step - loss: 0.4646 - accuracy: 0.7959 Epoch 3/100 250/250 [=====] - 2s 7ms/step - loss: 0.4301 - accuracy: 0.8083 Epoch 4/100 250/250 [=====] - 2s 6ms/step - loss: 0.4148 - accuracy: 0.8161 Epoch 5/100 250/250 [=====] - 2s 6ms/step - loss: 0.4001 - accuracy: 0.8248 Epoch 6/100 250/250 [=====] - 2s 7ms/step - loss: 0.3846 - accuracy: 0.8351 Epoch 7/100 250/250 [=====] - 2s 7ms/step - loss: 0.3722 - accuracy: 0.8426 Epoch 8/100 250/250 [=====] - 2s 7ms/step - loss: 0.3649 - accuracy: 0.8454 Epoch 9/100 250/250 [=====] - 2s 7ms/step - loss: 0.3615 - accuracy: 0.8468 Epoch 10/100 250/250 [=====] - 2s 7ms/step - loss: 0.3599 - accuracy: 0.8482 Epoch 11/100 250/250 [=====] - 2s 7ms/step - loss: 0.3582 - accuracy: 0.8489 Epoch 12/100 250/250 [=====] - 2s 7ms/step - loss: 0.3576 - accuracy: 0.8474 Epoch 13/100 250/250 [=====] - 2s 7ms/step - loss: 0.3565 - accuracy: 0.8495 Epoch 14/100 250/250 [=====] - 2s 6ms/step - loss: 0.3555 - accuracy: 0.8503 Epoch 15/100 250/250 [=====] - 2s 6ms/step - loss: 0.3556 - accuracy: 0.8511 Epoch 16/100 250/250 [=====] - 2s 6ms/step - loss: 0.3545 - accuracy: 0.8503 Epoch 17/100 250/250 [=====] - 2s 6ms/step - loss: 0.3534 - accuracy: 0.8514 Epoch 18/100 250/250 [=====] - 2s 7ms/step - loss: 0.3530 - accuracy: 0.8520 Epoch 19/100 250/250 [=====] - 2s 7ms/step - loss: 0.3516 - accuracy: 0.8526 Epoch 20/100 250/250 [=====] - 2s 7ms/step - loss: 0.3515 - accuracy: 0.8531 Epoch 21/100 250/250 [=====] - 2s 7ms/step - loss: 0.3506 - accuracy: 0.8543 Epoch 22/100 250/250 [=====] - 2s 7ms/step - loss: 0.3495 - accuracy: 0.8550 Epoch 23/100 250/250 [=====] - 2s 8ms/step - loss: 0.3489 - accuracy: 0.8556 Epoch 24/100 250/250 [=====] - 2s 6ms/step - loss: 0.3487 - accuracy: 0.8571 Epoch 25/100 250/250 [=====] - 2s 6ms/step - loss: 0.3478 - accuracy: 0.8566 Epoch 26/100 250/250 [=====] - 2s 7ms/step - loss: 0.3471 - accuracy: 0.8575 Epoch 27/100 250/250 [=====] - 2s 7ms/step - loss: 0.3470 - accuracy: 0.8580 Epoch 28/100 250/250 [=====] - 2s 7ms/step - loss: 0.3460 - accuracy: 0.8597 Epoch 29/100 250/250 [=====] - 2s 7ms/step - loss: 0.3457 - accuracy: 0.8586 Epoch 30/100 250/250 [=====] - 2s 7ms/step - loss: 0.3444 - accuracy: 0.8601 Epoch 31/100 250/250 [=====] - 2s 8ms/step - loss: 0.3442 - accuracy: 0.8596 Epoch 32/100 250/250 [=====] - 2s 8ms/step - loss: 0.3439 - accuracy: 0.8615 Epoch 33/100 250/250 [=====] - 2s 7ms/step - loss: 0.3430 - accuracy: 0.8583 Epoch 34/100 250/250 [=====] - 2s 8ms/step - loss: 0.3430 - accuracy: 0.8612 Epoch 35/100 250/250 [=====] - 2s 8ms/step - loss: 0.3429 - accuracy: 0.8583 Epoch 36/100 250/250 [=====] - 2s 8ms/step - loss: 0.3422 - accuracy: 0.8597 Epoch 37/100 250/250 [=====] - 2s 8ms/step - loss: 0.3416 - accuracy: 0.8593 Epoch 38/100 250/250 [=====] - 2s 8ms/step - loss: 0.3415 - accuracy: 0.8606 Epoch 39/100 250/250 [=====] - 2s 8ms/step - loss: 0.3407 - accuracy: 0.8612 Epoch 40/100 250/250 [=====] - 2s 9ms/step - loss: 0.3409 - accuracy: 0.8615 Epoch 41/100 250/250 [=====] - 2s 8ms/step - loss: 0.3407 - accuracy: 0.8596 Epoch 42/100 250/250 [=====] - 2s 8ms/step - loss: 0.3403 - accuracy: 0.8627 Epoch 43/100 250/250 [=====] - 2s 8ms/step - loss: 0.3401 - accuracy: 0.8605 Epoch 44/100 250/250 [=====] - 2s 9ms/step - loss: 0.3401 - accuracy: 0.8622 Epoch 45/100 250/250 [=====] - 2s 8ms/step - loss: 0.3402 - accuracy: 0.8594 Epoch 46/100 250/250 [=====] - 2s 8ms/step - loss: 0.3403 - accuracy: 0.8597 Epoch 47/100 250/250 [=====] - 2s 7ms/step - loss: 0.3390 - accuracy: 0.8609 Epoch 48/100 250/250 [=====] - 2s 7ms/step - loss: 0.3395 - accuracy: 0.8600 Epoch 49/100 250/250 [=====] - 2s 7ms/step - loss: 0.3391 - accuracy: 0.8595 Epoch 50/100 250/250 [=====] - 2s 7ms/step - loss: 0.3393 - accuracy: 0.8599 Epoch 51/100 250/250 [=====] - 2s 8ms/step - loss: 0.3390 - accuracy: 0.8600 Epoch 52/100 250/250 [=====] - 2s 7ms/step - loss: 0.3385 - accuracy: 0.8616 Epoch 53/100 250/250 [=====] - 2s 7ms/step - loss: 0.3390 - accuracy: 0.8615 Epoch 54/100 250/250 [=====] - 2s 8ms/step - loss: 0.3388 - accuracy: 0.8600 Epoch 55/100 250/250 [=====] - 2s 7ms/step - loss: 0.3387 - accuracy: 0.8593 Epoch 56/100 250/250 [=====] - 2s 7ms/step - loss: 0.3379 - accuracy: 0.8616 Epoch 57/100 250/250 [=====] - 2s 7ms/step - loss: 0.3382 - accuracy: 0.8627 Epoch 58/100 250/250 [=====] - 2s 7ms/step - loss: 0.3378 - accuracy: 0.8619 Epoch 59/100 250/250 [=====] - 2s 6ms/step - loss: 0.3380 - accuracy: 0.8618 Epoch 60/100 250/250 [=====] - 1s 6ms/step - loss: 0.3376 - accuracy: 0.8627 Epoch 61/100 250/250 [=====] - 2s 6ms/step - loss: 0.3385 - accuracy: 0.8605 Epoch 62/100 250/250 [=====] - 2s 7ms/step - loss: 0.3376 - accuracy: 0.8625 Epoch 63/100 250/250 [=====] - 2s 7ms/step - loss: 0.3384 - accuracy: 0.8620 Epoch 64/100 250/250 [=====] - 2s 7ms/step - loss: 0.3377 - accuracy: 0.8624 Epoch 65/100 250/250 [=====] - 2s 7ms/step - loss: 0.3377 - accuracy: 0.8636 Epoch 66/100 250/250 [=====] - 2s 7ms/step - loss: 0.3375 - accuracy: 0.8625 Epoch 67/100 250/250 [=====] - 2s 7ms/step - loss: 0.3375 - accuracy: 0.8610 Epoch 68/100 250/250 [=====] - 2s 7ms/step - loss: 0.3374 - accuracy: 0.8635 Epoch 69/100 250/250 [=====] - 2s 7ms/step - loss: 0.3371 - accuracy: 0.8640 Epoch 70/100 250/250 [=====] - 2s 7ms/step - loss: 0.3368 - accuracy: 0.8644 Epoch 71/100 250/250 [=====] - 2s 7ms/step - loss: 0.3375 - accuracy: 0.8639 Epoch 72/100 250/250 [=====] - 2s 7ms/step - loss: 0.3369 - accuracy: 0.8626 Epoch 73/100 250/250 [=====] - 2s 7ms/step - loss: 0.3372 - accuracy: 0.8639 Epoch 74/100 250/250 [=====] - 2s 6ms/step - loss: 0.3368 - accuracy: 0.8643 Epoch 75/100 250/250 [=====] - 2s 7ms/step - loss: 0.3372 - accuracy: 0.8618 Epoch 76/100 250/250 [=====] - 2s 7ms/step - loss: 0.3367 - accuracy: 0.8631 Epoch 77/100 250/250 [=====] - 2s 7ms/step - loss: 0.3367 - accuracy: 0.8636 Epoch 78/100 250/250 [=====] - 2s 7ms/step - loss: 0.3360 - accuracy: 0.8636 Epoch 79/100 250/250 [=====] - 2s 7ms/step - loss: 0.3367 - accuracy: 0.8622 Epoch 80/100 250/250 [=====] - 2s 7ms/step - loss: 0.3366 - accuracy: 0.8630 Epoch 81/100 250/250 [=====] - 2s 6ms/step - loss: 0.3362 - accuracy: 0.8644 Epoch 82/100 250/250 [=====] - 2s 6ms/step - loss: 0.3363 - accuracy: 0.8633 Epoch 83/100 250/250 [=====] - 2s 6ms/step - loss: 0.3358 - accuracy: 0.8631 Epoch 84/100 250/250 [=====] - 2s 6ms/step - loss: 0.3359 - accuracy: 0.8633 Epoch 85/100 250/250 [=====] - 2s 7ms/step - loss: 0.3359 - accuracy: 0.8656 Epoch 86/100 250/250 [=====] - 2s 7ms/step - loss: 0.3361 - accuracy: 0.8624 Epoch 87/100 250/250 [=====] - 2s 7ms/step - loss: 0.3363 - accuracy: 0.8646 Epoch 88/100 250/250 [=====] - 2s 7ms/step - loss: 0.3355 - accuracy: 0.8656 Epoch 89/100 250/250 [=====] - 2s 7ms/step - loss: 0.3361 - accuracy: 0.8625 Epoch 90/100 250/250 [=====] - 2s 7ms/step - loss: 0.3359 - accuracy: 0.8621 Epoch 91/100 250/250 [=====] - 2s 7ms/step - loss: 0.3351 - accuracy: 0.8641 Epoch 92/100 250/250 [=====] - 2s 7ms/step - loss: 0.3355 - accuracy: 0.8641 Epoch 93/100 250/250 [=====] - 2s 7ms/step - loss: 0.3353 - accuracy: 0.8620 Epoch 94/100 250/250 [=====] - 2s 7ms/step - loss: 0.3356 - accuracy: 0.8654 Epoch 95/100 250/250 [=====] - 2s 7ms/step - loss: 0.3354 - accuracy: 0.8640 Epoch 96/100 250/250 [=====] - 2s 8ms/step - loss: 0.3345 - accuracy: 0.8635 Epoch 97/100 250/250 [=====] - 2s 8ms/step - loss: 0.3355 - accuracy: 0.8648 Epoch 98/100 250/250 [=====] - 2s 7ms/step - loss: 0.3347 - accuracy: 0.8640 Epoch 99/100 250/250 [=====] - 2s 7ms/step - loss: 0.3350 - accuracy: 0.8655 Epoch 100/100 250/250 [=====] - 2s 7ms/step - loss: 0.3349 - accuracy: 0.8635 -keras.callbacks.History at 0x26bac98edf0></div>
----------	--

Part 4 - Making the predictions and evaluating the model

Predicting the result of a single observation

Homework

Use our ANN model to predict if the customer with the following informations will leave the bank:

Geography: France

Credit Score: 600

Gender: Male

Age: 40 years old

Tenure: 3 years

Balance: \$ 60000

Number of Products: 2

Does this customer have a credit card? Yes

Is this customer an Active Member: Yes

Estimated Salary: \$ 50000

So, should we say goodbye to that customer?

Solution

In [25]:	<pre>print(ann.predict(sc.transform([[1, 0, 0, 600, 1, 40, 3, 60000, 2, 1, 1, 50000]])) > 0.5) #predicting desired train data #False means value is close to zero-->means customer will stay according to target dataset</pre>
----------	--

[[False]]

Therefore, our ANN model predicts that this customer stays in the bank!

Important note 1: Notice that the values of the features were all input in a double pair of square brackets. That's because the "predict" method always expects a 2D array as the format of its inputs. And putting our values into a double pair of square brackets makes the input exactly a 2D array.

Important note 2: Notice also that the "France" country was not input as a string in the last column but as "1, 0, 0" in the first three columns. That's because of course the predict method expects the one-hot-encoded values of the state, and as we see in the first row of the matrix of features X, "France" was encoded as "1, 0, 0". And be careful to include these values in the first three columns, because the dummy variables are always created in the first columns.

Predicting the Test set results

In [26]:	<pre>y_pred = ann.predict(X_test) y_pred = (y_pred > 0.5) print(np.concatenate((y_pred.reshape(len(y_pred),1), y_test.reshape(len(y_test),1)),1)) #Comparing between predicting result(0 column) vs real result(1 column)</pre>
----------	--

```
[[0 0]
 [0 1]
 [0 0]
 ...
 [0 0]
 [0 0]
 [0 1]]
```

Making the Confusion Matrix

In [27]:	<pre>from sklearn.metrics import confusion_matrix, accuracy_score cm = confusion_matrix(y_test, y_pred) print(cm) accuracy_score(y_test, y_pred) #Confusion matrix-->Finding accuracy #1516 correct prediction-Customer stays in the bank,205 leaves from the bank & 79 inaccurate prediction-Customer leaves from the bank,200 -Customer</pre>
----------	--

```
[[1526  69]
 [ 205 200]]
0.863
```

In []:	
---------	--